

# Blobs and SIFT features

---

Morten R. Hannemose, [mohan@dtu.dk](mailto:mohan@dtu.dk)

March 25, 2022

02504 Computer vision course lectures,  
DTU Compute, Kgs. Lyngby 2800, Denmark



# Learning objectives



After this lecture you should be able to:

- implement and use blob detection using Difference-of-Gaussians
- analyse and use SIFT features and feature matching

# Similarity

## Pixel-wise comparison

- Shift of single pixel vs. two views



# Similarity

## Basic idea

- Locally appearance between views is the same
- Variation can be handled via invariances



# Local image features

## SIFT – key elements

- Features localized at interest points
- Adapted to scale and invariant to appearance changes



# SIFT – scale invariant feature transform (Lowe, 2004)

- Scale-space blob detection – difference of Gaussians
- Interest point localization
- Orientation assignment
- Interest point descriptor
- Note – SIFT is one example of interest point feature

# Harris corners and BLOBs

Harris corners are features that have a large change of intensity in two orthogonal directions. They are:

- local,
- can be found at different scales by changing the Gaussian filters, and
- work in rotated frames.

Harris corners are found by first order derivatives whereas blobs are response to second order image derivatives.

# BLOBs – Binary Large OBjects

Correspond to:

- a dark area surrounded by brighter intensities or,
- a bright area surrounded by darker intensities



# Hessian

The Hessian matrix contains the second order derivatives

$$\mathbf{H}(x, y) = \begin{bmatrix} I_{xx}(x, y) & I_{xy}(x, y) \\ I_{xy}(x, y) & I_{yy}(x, y) \end{bmatrix},$$

where

$$I_{xx}(x, y) = \frac{\partial^2 I(x, y)}{\partial x^2}, \quad I_{yy}(x, y) = \frac{\partial^2 I(x, y)}{\partial y^2}, \quad \text{and} \quad I_{xy}(x, y) = \frac{\partial^2 I(x, y)}{\partial x \partial y}.$$

# Curvature

Second order derivatives measure curvature.

Eigenvalues ( $\lambda_1, \lambda_2$ ) measure the principal curvature, i.e. degree of change in intensity

Eigenvectors measure direction of change where the eigenvector corresponding to the largest eigenvalue ( $\lambda_1$ ) is the direction of most change and the second is orthogonal to that.

# Blob detection with Hessian

Similar to the Harris corner detector, we can use either of the measures

$$\det(\mathbf{H}) = \lambda_1 \lambda_2,$$


$$\text{trace}(\mathbf{H}) = \lambda_1 + \lambda_2,$$

where  $\lambda_i$  are the eigenvalues of the Hessian.


$\det(\mathbf{H})$  is the Gaussian curvature.

$\text{trace}(\mathbf{H}) = \nabla^2 I$  is the Laplacian, which we use for blob detection.


# The Laplacian



# Two Gaussians with different variance



# Difference of Gaussians vs Laplacian



# Blob detection with DoG

An approximation of  $\text{trace}(\mathbf{H})$  is the Difference-of-Gaussian (DoG).

Consider the following convolution with two different Gaussian kernels:

$$\nabla^2 I \approx D_\sigma = (G_{k\sigma} - G_\sigma) * I = L_{k\sigma} - L_\sigma,$$

where  $G_\sigma$  is a Gaussian convolution with a width of  $\sigma$  and  $k > 1$  is a scale factor.

## Blob detection with DoG



Not that kind of dog..



# SIFT – Scale invariance



# SIFT – Scale invariance

- Using difference of Gaussians for blob detection

$$\begin{aligned} D(x, y, \sigma) &= ((G_{k\sigma} - G_{\sigma}) * I)(x, y) \\ &= L(x, y, k\sigma) - L(x, y, \sigma) \end{aligned}$$

# SIFT

We need to ensure that the DoG is scale normalized (the DoG kernel)

$$\sigma^2 \nabla^2 G$$

If we take an offset in the heat equation, we have

$$\sigma \nabla^2 G = \frac{\partial G}{\partial \sigma} \approx \frac{G_{k\sigma} - G_{\sigma}}{k\sigma - \sigma}$$

# SIFT

Multiplying both sides with  $\sigma$  we get

$$\begin{aligned}\sigma^2 \nabla^2 G &\approx \frac{\sigma}{k\sigma - \sigma} (G_{k\sigma} - G_\sigma) \\ &= \frac{1}{k - 1} (G_{k\sigma} - G_\sigma)\end{aligned}$$

From this we get

$$(k - 1)\sigma^2 \nabla^2 G \approx G_{k\sigma} - G_\sigma$$

Since  $k$  is constant over scales it does not affect the relative response of the DoG kernel.

# Gaussian scale space – Efficient


- Convolution of two Gaussians yield a new Gaussian
- Generate scale space by iteratively blurring already blurred images more
  - Requires smaller Gaussian kernels
- $k = 2^{\frac{1}{3}}$

# Gaussian scale space – Efficient


- Convolution of two Gaussians yield a new Gaussian
- Generate scale space by iteratively blurring already blurred images more
  - Requires smaller Gaussian kernels
- $k = 2^{\frac{1}{3}}$
- $\sigma$  doubles after three images, the image is downsampled.
  - This is an **octave**.

# SIFT – Estimation of DoG

## Difference of Gaussians



## Extrema localization



# SIFT – Interest point localization

- Taylor approximation to second degree of 2D surface

$$D(x) = D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$

- Setting the derivative of  $D(x)$  to zero

$$\hat{\mathbf{x}} = -\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2} \frac{\partial D}{\partial \mathbf{x}}$$



# SIFT – Interest point localization

- We get

$$D(\hat{\mathbf{x}}) = D + \frac{1}{2} \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}}$$

- $|D(\hat{\mathbf{x}})| > 0.03$  otherwise the point is discarded.

## SIFT – Interest point along edges discarded

- The eigenvalues of the Hessian are proportional to the principal curvatures

$$\mathbf{H} = \begin{bmatrix} I_{xx} & I_{xy} \\ I_{xy} & I_{yy} \end{bmatrix} = \begin{bmatrix} \frac{\partial^2 I}{\partial x^2} & \frac{\partial^2 I}{\partial x \partial y} \\ \frac{\partial^2 I}{\partial x \partial y} & \frac{\partial^2 I}{\partial y^2} \end{bmatrix}$$

## Interest point along edges discarded

- $\lambda_1$  and  $\lambda_2$  are the eigenvalues of the Hessian

$$\text{trace}(\mathbf{H}) = I_{xx} + I_{yy} = \alpha + \beta$$

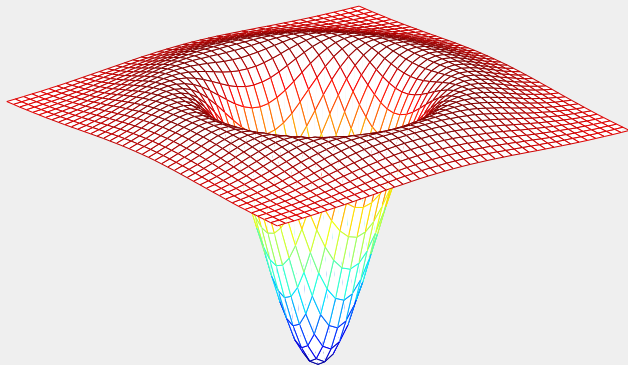
$$\det(\mathbf{H}) = I_{xx}I_{yy} - I_{xy}^2 = \alpha\beta$$

Points are kept if

$$\frac{\text{trace}(\mathbf{H})^2}{\det(\mathbf{H})} < \frac{(r+1)^2}{r},$$

where  $r = 10$  (found to be a good heuristic)

# DoG measure




Features have  $|\text{DoG}(I)| > \tau$ , where  $\tau$  is a threshold.

# DoG in scale pyramids

Scale pyramids are smaller scaled versions of the same image.

Scaled DoG is subtraction between all layers.



# Scale space blobs and DoG

DoG at different scales makes for a **scale invariant feature detector**.

Small and large details are recoverable in different DoGs.



## SIFT – Orientation assignment

Compute the orientation of gradients in a small region around the BLOB.

$$m(x, y) = \sqrt{L_x^2 + L_y^2}$$

$$\theta(x, y) = \arctan 2(L_y, L_x)$$

Where

$$L_x = L(x + 1, y) - L(x - 1, y)$$

$$L_y = L(x, y + 1) - L(x, y - 1)$$

# SIFT – Orientation assignment

- Compute circular histogram of gradient orientations
  - Weighted by magnitude, smoothed, and has 36 bins



# SIFT – Orientation assignment

- Compute circular histogram of gradient orientations
  - Weighted by magnitude, smoothed, and has 36 bins
- Use peak in histogram to assign orientation of point
- This introduces **rotation invariance**.
- Can we have multiple peaks in histogram?

# SIFT – Orientation assignment


- Compute circular histogram of gradient orientations
  - Weighted by magnitude, smoothed, and has 36 bins
- Use peak in histogram to assign orientation of point
- This introduces **rotation invariance**.
- Can we have multiple peaks in histogram?
  - Yes, this can happen at e.g. corners.
  - Create a new point at the same location if peak is over 80% of max.

# SIFT – Invariances

- Position
- Scale
- Rotation
- Linear intensity change
- Perspective changes?



# SIFT – Descriptor

- Create local patch at scale and orientation of point
- Build a histogram of local gradient orientations




- Normalized using  $L_2$  norm:  $\mathbf{d}_n = \frac{1}{\sqrt{\sum_{i=1}^{128} \mathbf{d}(i)^2}} \mathbf{d}$

# Taylor SIFT



# Taylor SIFT



# SIFT – Matching of descriptors

- Use Euclidean distance between normalized vectors

$$\delta(\mathbf{d}_i, \mathbf{d}_j) = \sqrt{\sum_{n=1}^{128} (d_{i,n} - d_{j,n})^2}$$

- Note – for comparison the square root is not needed

# RootSIFT

Simple trick to improve SIFT matching

- SIFT is a histogram
- Euclidean distance is dominated by large values
- RootSIFT is a transformation that measures distance using the Hellinger kernel.
  - L1 normalize
  - Take the square root of each element
  - L2 normalize the resulting vector
- Compare using Euclidean distance



# SIFT – Matching of descriptors

- For each feature in image 1 ( $d_{1,i}$ ) find the closest feature in image 2 ( $d_{2,j}$ )
  - This will give a lot of incorrect matches
- Cross checking
  - Only keep matches where  $d_{2,j}$  is also the closest to  $d_{1,i}$  of all features in image 1

# SIFT – Matching of descriptors

- For each feature in image 1 ( $d_{1,i}$ ) find the closest feature in image 2 ( $d_{2,j}$ )
  - This will give a lot of incorrect matches
- Cross checking
  - Only keep matches where  $d_{2,j}$  is also the closest to  $d_{1,i}$  of all features in image 1
- Ratio test
  - Compute the ratio between the closest and second closest match, and keep where this is below a threshold, e.g. 0.7.

# SIFT – Summary

- SIFT is both a feature detector and descriptor
- Find local extrema of DoGs in scale space
- Place patch oriented along local gradients
- Compute histograms of gradients.
- Allows matching of images invariant to: scale, rotation, illumination and viewpoint
- Partly visible objects can be matched

# Other descriptors

SIFT is widely used.

- ORB
- SURF
- BRISK
- FAST

# Learned descriptors

- Deep Learning has created improved feature descriptors.
- Mostly in improvement in invariance to changing lighting.
- Some examples:
  - R2D2: Repeatable and Reliable Detector and Descriptor  
<https://github.com/naver/r2d2>
  - Superpoint  
<https://github.com/magicLeap/SuperPointPretrainedNetwork>

# Exercise

Build a blob detector and match points with SIFT detector.

Python: Recommended to use OpenCV (4.2.0 or newer)

Matlab: Recommended to use VLFeat

<https://www.vlfeat.org/overview/sift.html>

# Midterm Evaluation

**Exercise time!**